

An Open-Source Event Based SCADA System for the Power Grid

Trevor Aron

A project report submitted to the Johns Hopkins University in conformity
with the requirements for the degree of Master of Science in
Engineering

Advisors: Dr. Yair Amir, Tom Tantillo

May 2017

Acknowledgements

I would like to thank Marco Platania. Marco Platania really introduced me to what SCADA actually is. We came up with the scenario that I would end up building together after scouring the internet for youtube videos of HMI demos. Marco was also the first person to realize the limitations of PvBrowser and split the HMI and Data Acquisition portions, leading to the creation of the new SCADA architecture I am now presenting.

I would like to thank Tom Tantillo. Tom Tantillo really spearheaded turning this project into the full Spire system, which will be his PhD thesis. He worked with me tirelessly to get this made. Tom has an amazing ability to get things done, and done well, no matter how impossible it seems. He has served as one of my main mentors, helping me with everything from writing posters to drafting emails. Much of this work was on his suggestion.

I must also thank Yair Amir. He was my professor in Intermediate Programming, Distributed Systems, and Advanced Distributed System. It was he who made me interested in systems. It was also he who had the vision for intrusion tolerant SCADA and fortunately invited me to help. I truly admire his desire to protect the nations infrastructure for the good of society.

Finally I would like to thank Amy Babay. Amy Babay has helped me in countless ways and is always there to answer my questions. She also made significant work towards Spire .

Contents

1	Acknowledgements	3
2	Introduction	5
2.1	Abstract	5
2.2	Background	5
2.2.1	SCADA	5
2.3	Motivations	6
2.3.1	SCADA Security	7
2.3.2	Scalability	8
2.3.3	Replication	9
2.3.4	Adoption	9
3	Open Source Event-Based Architecture	11
3.1	Overview	11
3.2	pvbrowser HMI	12
3.3	Master	13
3.4	Proxy	13
3.5	Open PLC	15
4	Case Study: Power Distribution Scenario	17
4.1	Overview	17
4.2	HMI	17
4.3	Master	18
4.4	Proxy	18
4.5	PLC Emulation	18
	Bibliography	21

Introduction

2.1 Abstract

This document presents the construction of an open-source event-based SCADA architecture. This construction is used as a backbone of the Spire Intrusion Tolerant SCADA system. The aim of this work is to build a SCADA system that is scalable, and has an easy to work with architecture that is still backwards compatible with SCADA equipment. The key component of this architecture is the RTU/PLC Proxy. This additional machine allows the SCADA system to be fully backwards compatible with devices while making sure that updates from devices are delivered to the core system as changes occur. Additionally this extra component provides an additional layer of security to the SCADA system.

2.2 Background

2.2.1 SCADA

SCADA stands for Supervisory Control and Data Acquisition. They are systems used to monitor and control physical devices in critical infrastructure applications. Such applications include railways, electrical grids, power generation, and factories. [?].

SCADA systems are very different depending on the scenario being monitored and control. These specific conditions can influence the topology of the network and the protocols being used to communicate between the Master and PLC. [?] They are also mainly vendor locked, changing wildly depending on what company built the system. However, all SCADA systems share the basic architecture described in Figure 2.1. This shows the three main portions of any control system: the Human Machine Interface (HMI), the SCADA Master, and the Programmable Logic Controller (PLC) or Remote Terminal Unit (RTU).

The Human Machine Interface is the machine that visualizes the data to a Human operator, and allows said operator to make changes to the system. What an HMI looks like is very specific on the critical infrastructure system being monitored. Usually, they are set up with images that look like the actual system – an HMI monitoring water levels may have an image of a tank with a variable amount of water in it. They can also just be numbers on a screen – it really depends on the system and the company making the HMI.

Programmable Logic Controllers and Remote Terminal Units are the devices in the system actually being monitored and controlled. These machines are in the field and will either be integrated with the

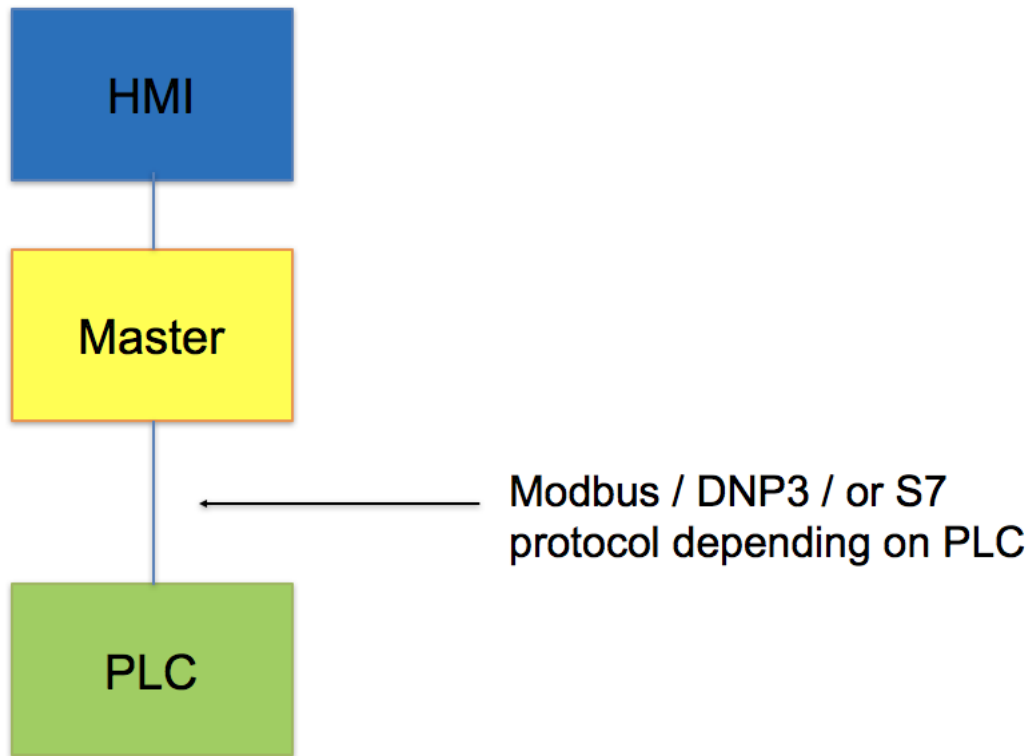


Figure 2.1: Architecture of a SCADA System

devices. RTU's are a little more intelligent than PLC's and are usually outfitted with advanced communication systems such as radio. They may also perform some small automated control tasks. They are specifically for SCADA systems. PLC's are a bit more generic, and most hobbyist boards (such as Arduino's) can be considered as a PLC. RTU's and PLC's in SCADA systems are usually vendor created, closed source devices. They also communicate with legacy communication protocols, which have been around since the seventies and eighties. There are many different protocols, three common ones for power grids are Modbus, DNP3, and Siemens S7. Beyond these protocols, the American Gas Association's AGA-12 standard states that there are between 150-200 different SCADA communication protocols. [?]

The SCADA Master is the brain of the system. All commands from the HMI first go to the SCADA Master before they are sent to a PLC. The SCADA Master typically speaks whatever protocol languages the PLC's and RTU's it monitors can speak. The SCADA Master typically has automated control and alerting abilities. They also contain a historian which keeps an audit trail of operational data. [?]

2.3 Motivations

There are four main motivations for building this open source event based architecture. The first is Security: the RTU Proxy component in the event based SCADA architecture provides a powerful layer of security. The second is scalability: an event based system pushes less data over the network. The third is

for replication: it is much easier to replicate an event based system than a polling based system. The final is the motivation for open source: adoption.

2.3.1 SCADA Security

SCADA systems were designed in the earlier days of computing before cyber security was a real consideration. However, recent evidence shows that SCADA systems need to make security a priority.

Stuxnet is a virus, first identified in 2010, that was discovered to target Siemens SCADA systems. Specifically, it was aimed at the Siemens SCADA systems that controlled nuclear enrichment facilities in Iran. The system was air gapped, but the virus entered by copying itself on remote drives and spreading in private networks. It affected computers that were used to modify and program PLC's and used this to insert malicious code into a PLC that would make it function incorrectly.

The onset of Stuxnet lead the world to look carefully at SCADA systems and their cyber security. The test in [?] shows the extent of the threat. They set up a honey pot of PLC's connected to the internet. In a month time frame there where over 39 attacks from 14 different countries. SCADA systems have traditionally used air gapped networks to create security, but many of these PLC's and RTU's are configured incorrectly or connected to the internet such that employees can access and modify them from home. One can find PLC's on the internet with a basic search. The website Shodan is a search engine designed to discover Internet of Things (IOT) devices on the internet. They have an entire section devoted to finding industrial control equipment by looking for IP connected devices with common SCADA protocol ports open. [?]

The weakest parts of the system are the devices. They are hard to harden: PLC's may have very limited computational abilities. They also often run on real time operating systems which are lighter weight, but provide less security. Thus, many PLC's can not support running a firewall. [?] In addition, because they run on real time operating systems, they are more susceptible to disruption from denial of service attacks. Beyond the PLC itself, the protocols they speak are not secure. Most typically do not support cryptography primitives. This is also due to the limited computing power of the devices. Thus most communication to the devices is unencrypted and most commands are unauthenticated. For instance, Modbus is sent over the wire completely unencrypted, with no integrity checks, and no authentication. If one is on the same networks as a device that speaks Modbus, they can do whatever they want. [?] In addition, because the world of SCADA is vendor locked, many of these protocols have been programmed from scratch for a particular device. Many of these implementations are not robust and have bugs. In our own experiments with the ASE Test Set 2000 RTU emulation device we found that it had bugs in it's implementation of DNP3. These bugs are often the entry way for intruders.

There have been many efforts to harden PLC's. For instance, the efforts in this paper [?] propose a new Modbus protocol that is translated by a middle gateway device into regular Modbus for backwards compatibility with devices. The secure modbus However, this breaks backwards compatibility with masters. A more broad based approach has also been to attach machines at both ends, one at the SCADA Master, and the other at the RTU or PLC. [?] This acts as a bump in the wire in which data is encrypted.

To solve the firewall issue, there has also been efforts to place small firewalls in front of each PLC in a network. [?]

The solution that I propose, the RTU proxy, solves these issues. The RTU Proxy is a machine that will be placed in front of RTU's or PLC's on the network. It translates commands from the Master into the specific PLC or RTU communication protocol for the corresponding device. This proxy runs a firewall for the devices that it speaks to, removing them from direct access to the wide area network. Since the proxy is a modern machine, and not a device board, it can also perform cryptographic primitives. Thus, the information it gets from the Master has authentication, integrity, and confidentiality. Since it speaks many differing SCADA protocols, it provides backwards compatibility for devices.

2.3.2 Scalability

There are two different SCADA protocol architectures. The first is a polling model, and the second is an event based model. An event based model consists of the device only sending updates to the master whenever there is a change of state. DNP3 is an event based protocol. The polling model consists of the master sending requests for updates to the devices at different polling intervals. Modbus is a polling protocol.

One of the issues with the polling model is that it doesn't scale. Each device takes a constant amount of bandwidth. This is expensive. Some devices may transmit a lot of data on each polling interval. There are some SCADA systems, like smart grids, are very large. There is an estimated 150,000,000 meters installed in Europe. [?]. If the devices only speak a protocol like Modbus, then there would be a massive amount of bandwidth used. In addition the historian would have to store much more data. Event based models are more efficient because only when a device's state has changes they send data.

The need for SCADA scalability has been recognized by others in the field. SAP and Schneider Electric have laid out their vision for the future in SCADA in the following paper: [?]. They address that it is impossible to perform polling in large scale SCADA systems, and suggest the transition to event driven infrastructures is the future. However, they stress that future SCADA systems must stay backwards compatible to work with existing devices. The authors of [?] were running a smart grid monitoring system. In order to scale this monitoring, they created a new protocol and communication pattern. The issue with this work, in greater adoption, is that it breaks backwards compatibility, and new devices need to be made. SCADA systems can use DNP3 devices and solve this issue, but not all devices speak DNP3.

The solution I propose has a device, the RTU proxy, colocated with devices. The RTU proxy speaks multiple SCADA protocols (currently Modbus and DNP3, but it can be extended to many more) and translates this information to a generic format the SCADA master understands. The Proxy can be designed such that it only pushes information when there is a change in state. This solution makes all SCADA systems have more homogenous communication patterns despite the devices that may need to be used.

2.3.3 Replication

One of the goals of this work was to make SCADA software such that the master could be replicated. Replication opens the door for making SCADA systems more available: a system could be replicated with Paxos to be fault tolerant, or a byzantine fault tolerant algorithm to make the system intrusion tolerant - allowing it to work even if components have been compromised by an intruder. Specifically, this system was used to make Spire, an intrusion tolerant SCADA system.

A former member of the DSN lab, John Kirsch, published a paper [?] about his efforts in making a Siemens SCADA system intrusion tolerant. He notes that intrusion tolerant replication systems, such as Prime, assume that updates are client driven (event based), while some SCADA systems process requests that are server driven (polling). This mismatch forced him to create the intrusion tolerant timeout protocol. Not only is this protocol non trivial to implement, it also requires additional Prime orderings which negatively impact system performance by increasing load.

By making an event based system, this extra intrusion tolerant timeout protocol is no longer needed, and backwards compatibility is still supported. This makes replication of the system much simpler, and conserves the number of updates that need to be ordered.

2.3.4 Adoption

Finally, the motivation to build this system entirely with open source components is to encourage usage and adoption in the SCADA community. The efforts of John Kirsch in [?] are unfortunately unavailable as Siemens decided not to release this product. By building a completely open source SCADA system, this work can be used as a base for SCADA security research to build off of. Spire is able to be released as open source because this system is uses all fully open source parts.

In addition, using open source components mean that this system will benefit as the differing parts get further research. OpenDNP3 [?] is currently maintained by a group that performs vulnerability research on SCADA protocols. We use this for our DNP3 implementation, and the security of our system will be improved from their work. We use OpenPLC [?] to emulate PLC's in this system. This integration benefits our system as Open PLC is being designed to be used as a vehicle for PLC security research, and when the security of Open PLC improves, so will our system. In addition Open PLC supports a wide variety of real PLC boards to deploy on, which may encourage others to adopt this solution.

Open Source Event-Based Architecture

3.1 Overview

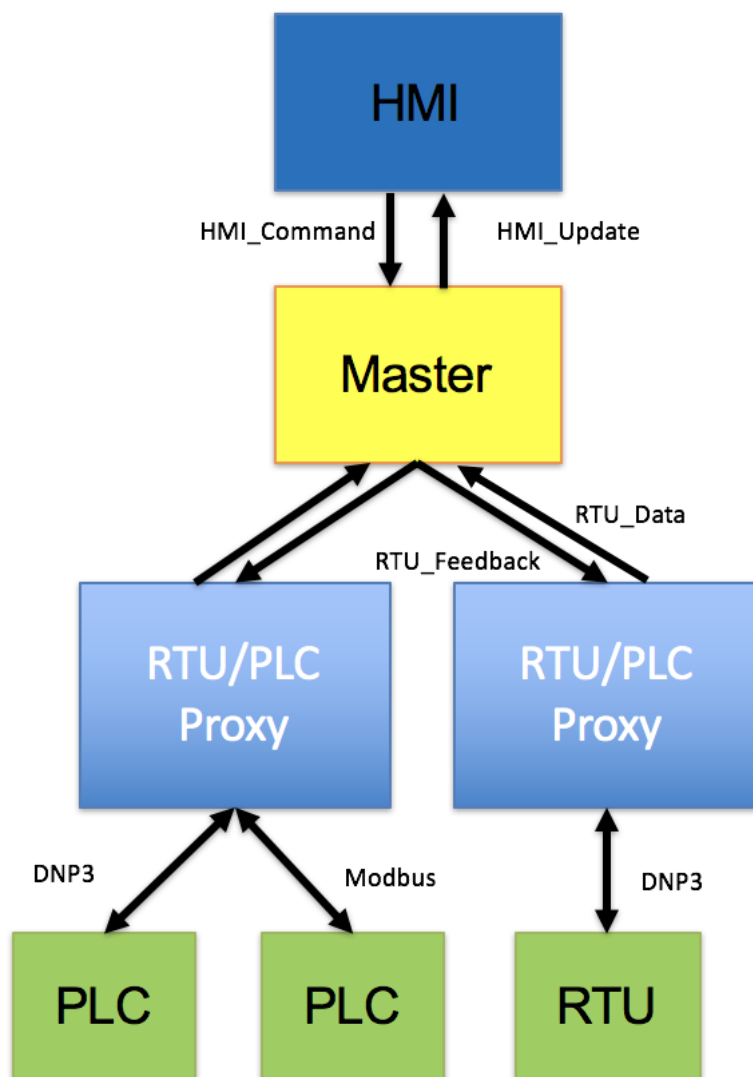


Figure 3.1: New Open Source Event Based Architecture

The architecture of the entire system is present in figure 3.1. The arrows describe the direction of message flow, with the labels being the type of messages. The system supports any configuration

of proxies and PLCs, with proxies able to support multiple PLC's of different protocol communication types. This configuration is described in the file `config/config.json`. Settings here are used to let all components know about the other elements in the system and allow the system to be configured in a number of ways. Information such as the IP address of the PLC's, the communication protocols they use, and the ports they are listening on are described in this configuration file. The system currently supports DNP3 and Modbus, but it is extendable to allow easy extension to future protocols. The cJSON library is used to support this functionality [?]

The packet types `HMI_Command`, `HMI_Update`, `RTU_Feedback`, and `RTU_Data` are described in the file `common/scada_packets.h`. The PLC Proxy gets data from it's PLC's or RTU's speaking either Modbus or DNP3, and sends this up to the master in a `RTU_Data` message whenever there is a change of state. This message is not in any SCADA protocol, but contains bundled information about the PLC's that the master needs. The master processes this, updates it's state, and then sends the HMI a `HMI_Command` message that has all the information the HMI needs to visualize the scenario. The HMI processes this to present the viewer with the monitoring information. If the monitor clicks a button, the HMI recognizes this event and sends a `HMI_Command` message to the master. The master gets this, uses the config file to figure out what proxy to forward this to, and sends a `RTU_Feedback` message to the proxy. The proxy uses the config to figure out what PLC to route this message to, and translates the feedback message into a command message in the given PLC's communication protocol.

3.2 pvbrowser HMI

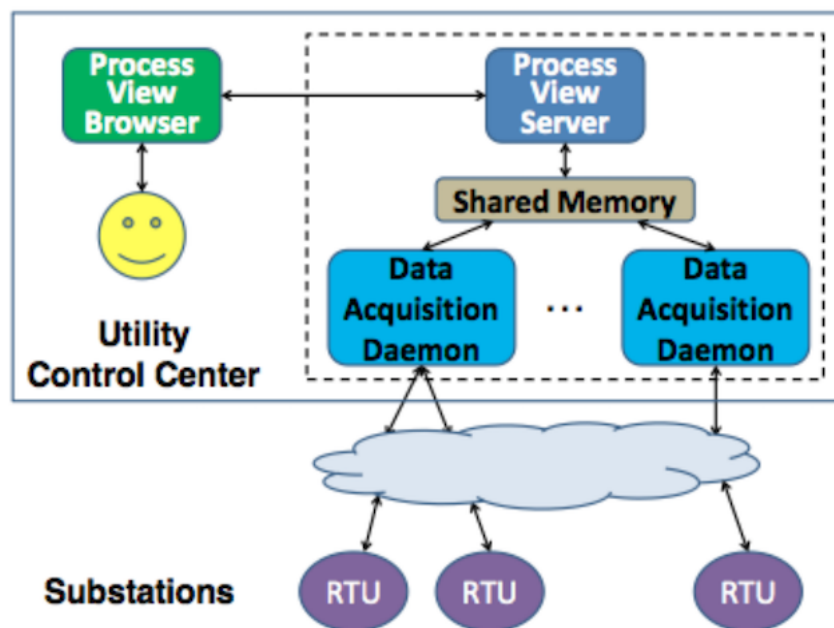


Figure 3.2: PvBrowser Stock Architecture

We use pvbrowser ([?]) as our HMI. pvbrowser is an open source SCADA software suite. It is used to manage real SCADA deployments, such as the Romanian power grid. It is a full SCADA solution - it provides an HMI and a master with Modbus data acquisition daemons that can communicate with RTU's. It's architecture is pictured in 3.2. Early work made it clear that replicating pvbrowser has some of the same difficulties Kirsch ran into in [?] , but the software had a lot of great components, such as the HMI and Modbus communication, that we wished to keep. We changed pvbrowser so that the Process View Browser Component and Process View Server component were colocated on the same machine - this is now the HMI box. We then removed the shared memory and data acquisition daemons, and replaced them with a thread that communicates with the master, and supplies the pvbrowser thread with the most up to date information of how to visualize the system. When there is a button click, the pvbrowser thread sends a message to the master.

This code is located in hmi. hmi/master_exec.cpp contains the thread that reads from the master and updates the pvbrowser data structures. hmi/mask1_slots.h contains the code both to visualize the HMI based on the current data structures, as well as the code to send HMI_Command messages when buttons are clicked.

3.3 Master

The SCADA master is built from scratch in C. It is located in master/scada_master.c and its data structures are in master/structs.h. The master uses the config file to figure out what RTU's or PLC's to route requests to. The master server runs a switch statement waiting for different message types. When it receives a HMI_Command message, it calls the read_from_hmi method which discovers what the corresponding RTU_Feedback message should be doing (what field of a RTU it will be modifying) and where to route it, and then sends it to the proxy. When it receives a RTU_Data message, it updates its data structures, then calls the process method. This method is unique to each individual scenario. Afterwards it will craft a HMI_Update message and send it to the HMI.

3.4 Proxy

The architecture of the PLC Proxy is shown in figure 3.3. The Proxy process is located in proxy/proxy.c. When it runs, it checks its config file to see what PLC's or RTU's it's responsible for, and what processes they run. If there are any PLC's that speak Modbus, it spawns the process modbus/modbus_master. If it is responsible for any PLC's that speak DNP3 it spawns the process dnp3/dnp3_master. It is set up that if there are any more protocols implemented, it would spawn their master as well. The master processes are designed such that they can handle multiple PLC's of the same protocol. For example, if the proxy is responsible for three PLC's that speak Modbus, it will only spawn one Modbus master process. An IPC channel is created for each of the child data acquisition daemons such that the proxy can communicate to the child and the child can communicate with the proxy. When the proxy gets an RTU_Feedback message, it checks to see what protocol the destined PLC speaks, and forwards the RTU_Feedback message to the designated daemon. When it receives RTU_Data messages from its children, it forwards




Figure 3.3: Proxy Architecture

those messages to the SCADA master.

The Modbus process is a modified version of pvbrowsers's Modbus data acquisition daemon from pvbaddons [?]. Its code is located at `modbus/modbus_master.cpp`. It reads the config file and uses the information there to figure out where the PLC's that it is responsible for are located on the network. It then creates TCP connections with them and starts the <odbus polling protocol. Every time out it polls the PLC's at the location specified in the config file, and then forwards a `RTU_Data` message with the corresponding information to the proxy over IPC. If the Modbus master receives a `RTU_Feedback` message, it will translate this into the proper Modbus control message and forward this to the corresponding PLC.

The DNP3 master process uses the OpenDNP3 library ([?]) to implement DNP3 communication with its devices. DNP3 is a more advanced event based communication protocol that is common in power grid networks. OpenDNP3 provides a modern, C++11 programming API for implementing DNP3 communication. The code for the DNP3 daemon is in `dnp3/`. The file `dnp3/main.cpp` is responsible for reading the config file and recognizing what PLC's it has to communicate with, and starting a DNP3 session with those PLC's. It sets up call-back's for these PLC's, such that when they send an event

update, the code in `dnp3/callback.cpp` runs, and creates a `RTU_Data` message to be sent to the proxy process. The main thread also sets up IPC communication with the proxy such that when it gets a `RTU_Feedback` message it translates this into a DNP3 control message and sends it along to the corresponding PLC.

3.5 Open PLC

We use OpenPLC [?] to emulate PLC's and RTU's. It is very useful, as it allows us to emulate realistic PLC's that a SCADA system would have to control and monitor. In addition to emulation, OpenPLC software can be deployed to many hobby PLC boards to actually control devices.

OpenPLC is configured by creating a Ladder Logic (LD) or Structured Text (ST) description of the PLC. Variables can be mapped to register positions on the PLC, and manipulated with the Ladder Logic or Structured Text. These same registers map to Modbus registers. Devices can then communicate with the Open PLC via Modbus by polling for the desired registers. we extended OpenPLC to also map the registers to DNP3 addresses and communicate with masters via DNP3. This is done with OpenDNP3, and allows to emulate a wider gambit of devices.

Case Study: Power Distribution Scenario

4.1 Overview

To demonstrate this architecture used in a real SCADA system we built a Power Distribution scenario. The scenario is designed with ten PLC's providing the system with power switching information. An operator can see this, and use it to route power around different substation.

4.2 HMI

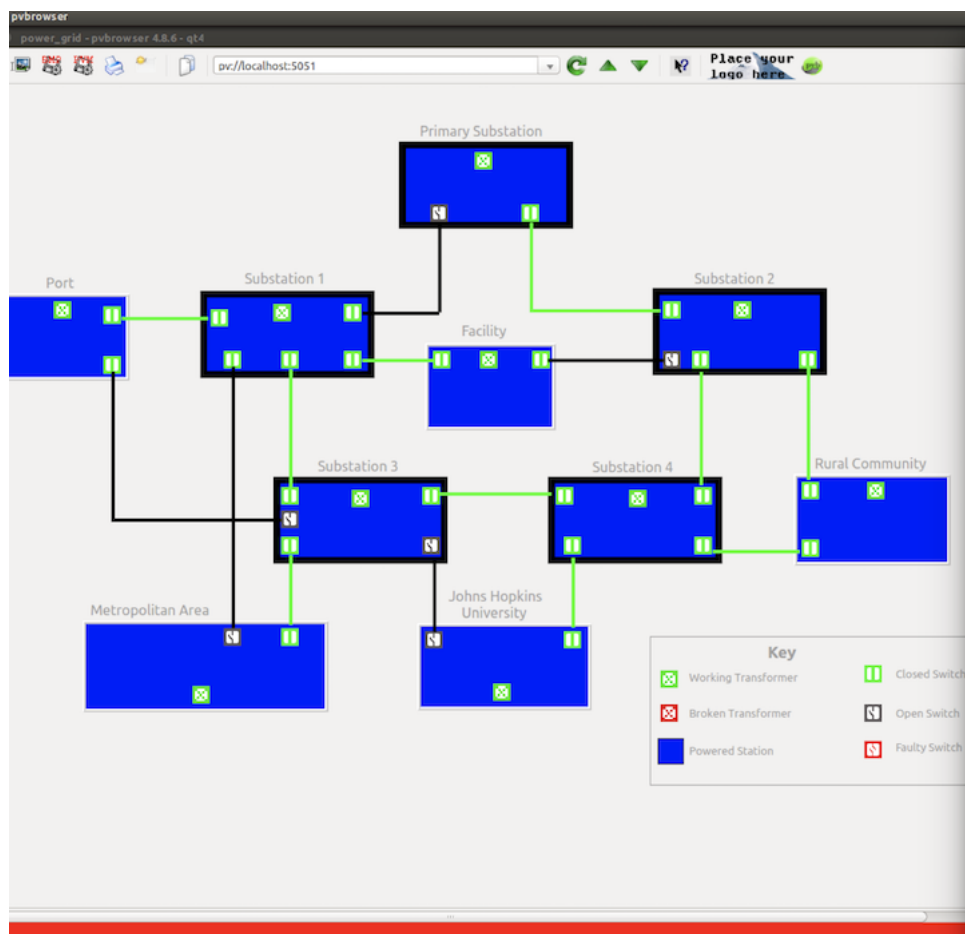


Figure 4.1: pvbrowser HMI

The HMI is pictured in figure 4.1. The boxes are substations - if they are blue there is power located at the substation. Power is generated at the primary substation. If a substation is black than that substation has no power. The 'x' in the middle of a substation is a transformer. If it is red, then no power can be distributed from that particular substation. If switches are green, they are closed, if they are black they are open, and if they are red they are tripped, which means they are broken and can't be modified before they are fixed. When switches are open or tripped, power can not flow through a line and thus it is black. When switches are both closed, and one of the substations has power, then power will flow through the line and it will be green.

To make changes to the system, the operator can click on either the transformers, or the switches, and the click will cause the system to change that PLC's state.

4.3 Master

The SCADA master for this scenario runs a version of Breadth First Search to determine what substations are powered, and what lines are carrying electricity. It does this because it knows from the PLC's only which switches are open or closed, and which transformers are on or off. Only the primary substation is known to be powered, and the rest of the information for the operator has to be extrapolated from the data the PLC's are providing. It takes the result of this process, puts it in an array, and sends it to the HMI in a `HMI_Update` message.

The master is able to translate `HMI_Command` messages, which come from the HMI and specify what item has been pressed, into `RTU_Feedback` messages, that are tagged to a specific PLC and Proxy, and, if the item pressed is a switch, contain the specific register number the switch should correspond to.

4.4 Proxy

In this system setup there is a proxy for each PLC, so a total of 10 proxies are run. This is because each PLC is supposed to be located at a different location.

The proxies protocol daemons knows how to translate a `RTU_Feedback` message into a corresponding Modbus or DNP3 message to send to the PLC or RTU. For DNP3, switch controls are Analog Output Commands and the transformers are CROB instructions. For Modbus, the switch controls send a set register command and the transformers are force coil commands.

4.5 PLC Emulation

To run this scenario we emulated 10 PLC's - one for each substation. The first three PLC's are emulated with OpenPLC using DNP3. The next two PLC's are emulated with OpenPLC using Modbus. To demonstrate our system's backwards compatibility, we also emulate five PLC's with the ASE Test Set

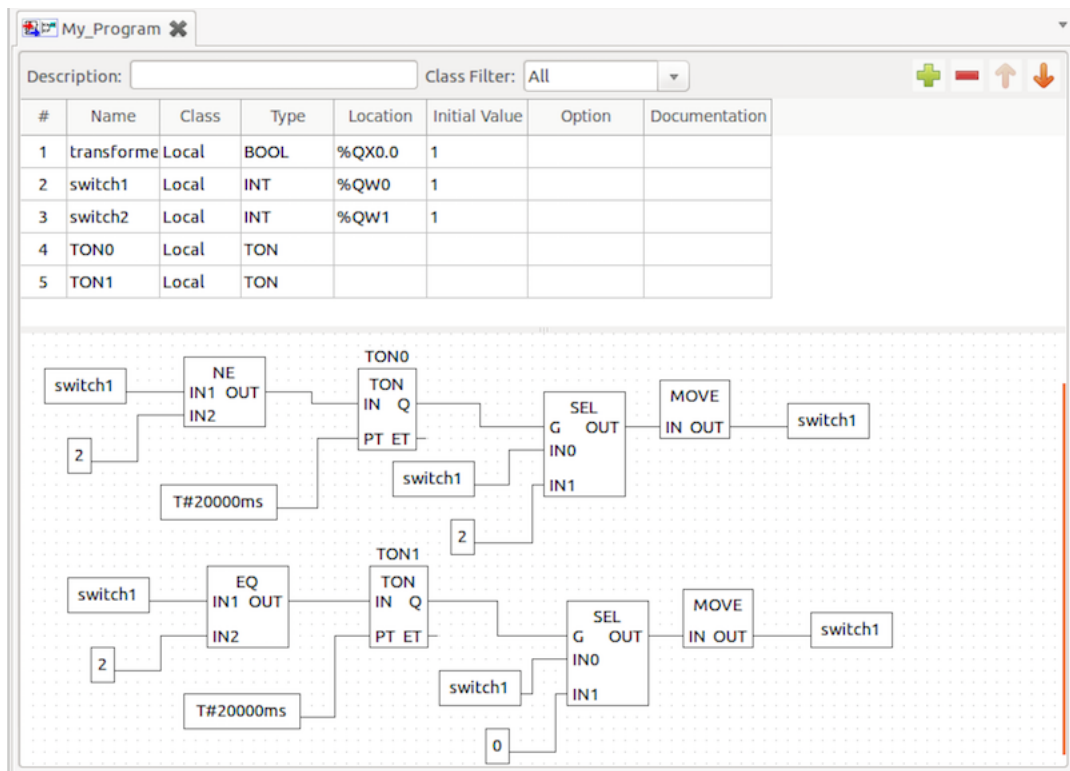


Figure 4.2: Ladder Logic for RTU 0 in Open PLC

2000 Device [?]. It can emulate Modbus or DNP3 RTU's and is used by Siemens to test their SCADA equipment. The next three PLCs are emulated with the ASE device through DNP3, and the last two are emulated with the ASE device through Modbus.

The PLC's contain data for the transformers and switches. A switch value of 0 is open, 1 is closed, and 2 is tripped. A transformer value of 1 is on and 0 is off. This data is represented in two different ways depending on if the device is Modbus or DNP3. If the device is Modbus, the transformer is a coil status at register zero, and the switches are holding registers at the register it's switch number should be (if a substation has three switches it will have holding registers at 0, 1, 2 to store that switch's data). For DNP3, the registers for the transformers and switches are at the same location, but instead of Coil Status and Holding Registers it is Binary Output Status and Analog Output Status.

In this scenario, PLC 0 (the primary substation) has been written with Ladder Logic in Open PLC. The associated LD program is shown in figure 4.2. This program trips the first switch every 20 seconds, and then un=trips it and sets the status to open every other 20 seconds.

Bibliography